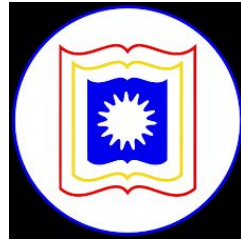


CSE4261: Neural Network and Deep Learning

Lecture: 21.05.2025



Sangeeta Biswas, Ph.D.
Associate Professor,
University of Rajshahi, Rajshahi-6205, Bangladesh

Regression Vs. Classification

Regression is a problem of predicting a continuous numeric value (e.g., a price, a temperature, a score) based on input data.

Regression



What will be the temperature tomorrow?




Fahrenheit

Classification



Will it be hot or cold tomorrow?



Fahrenheit

Computer Vision Tasks

Classification




CAT

**Classification
+ Localization**



CAT

Object Detection



CAT, DOG, DUCK

**Instance
Segmentation**




CAT, DOG, DUCK

Single object


Multiple objects

Classification

- Assigning objects to some pre-existing classes / categories / labels / groups.



Number of classes = 2




Number of classes = 4

What will We Use

- Linux Environment
- Tensorflow, Keras, PyTorch
- Python
- OpenCV
- Other Python Libraries such as matplotlib, scikit, pandas

MNIST Digit Dataset

- 70000 images 28 x 28
 - 60000 for training
 - 10000 for testing
- 10 classes
 - 0
 - 1
 - 2
 - 3
 - 4
 - 5
 - 6
 - 7
 - 8
 - 9



MNIST Fashion Dataset

- 70000 images 28 x 28
 - 60000 for training
 - 10000 for testing
- 10 classes
 - 0: T-shirt/top
 - 1: Trouser
 - 2: Pullover
 - 3: Dress
 - 4: Coat
 - 5: Sandal
 - 6: Shirt
 - 7: Sneaker
 - 8: Bag
 - 9: Ankle boot



Other Datasets for Classification

- CIFAR 10 dataset & CIFAR 100 dataset
 - <https://www.cs.toronto.edu/~kriz/cifar.html>
 - 60000, 32x32 colour images in 10 classes, with 6000 images per class
 - 50000 training images and 10000 test images
- CIFAR 100 dataset
 - <https://www.cs.toronto.edu/~kriz/cifar.html>
 - 100 classes containing 600 images each
 - 500 training images and 100 testing images per class
- ImageNet 1000 dataset
 - <https://www.image-net.org/>
 - 1000 object classes
 - 1281167 training images, 50000 validation images and 100000 test images.

Pre-trained Classifiers Provided by Keras API Team

- Models were trained by ImageNet 1K dataset for 1000 classes

<https://keras.io/api/applications/>

- Most of the time, we depend on these pre-trained models

Model	Size (MB)	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth	Time (ms) per inference step (CPU)	Time (ms) per inference step (GPU)
Xception	88	79.0%	94.5%	22.9M	81	109.4	8.1
VGG16	528	71.3%	90.1%	138.4M	16	69.5	4.2
VGG19	549	71.3%	90.0%	143.7M	19	84.8	4.4
ResNet50	98	74.9%	92.1%	25.6M	107	58.2	4.6
ResNet50V2	98	76.0%	93.0%	25.6M	103	45.6	4.4
ResNet101	171	76.4%	92.8%	44.7M	209	89.6	5.2
ResNet101V2	171	77.2%	93.8%	44.7M	205	72.7	5.4
ResNet152	232	76.6%	93.1%	60.4M	311	127.4	6.5
ResNet152V2	232	78.0%	94.2%	60.4M	307	107.5	6.6
InceptionV3	92	77.9%	93.7%	23.9M	189	42.2	6.9
InceptionResNetV2	215	80.3%	95.3%	55.9M	449	130.2	10.0
MobileNet	16	70.4%	89.5%	4.3M	55	22.6	3.4
MobileNetV2	14	71.3%	90.1%	3.5M	105	25.9	3.8
DenseNet121	33	75.0%	92.3%	8.1M	242	77.1	5.4
DenseNet169	57	76.2%	93.2%	14.3M	338	96.4	6.3

Code for Loading MNIST Fashion Dataset

```
#-- Load data
```

```
from tensorflow.keras.datasets import fashion_mnist
```

```
(trainX, trainY), (testX, testY) = fashion_mnist.load_data()
```

```
print(trainX.shape, trainX.dtype, trainY.shape, trainY.dtype)
```

```
#--- Cross check
```

```
plt.imshow(trainX[0])
```

```
plt.title(trainY[0])
```

```
plt.show()
```

```
plt.close()
```

Code for Loading Pre-trained Model

#--- Load the pre-trained model, VGG16, with head

```
from tensorflow.keras.applications import vgg16
```

```
vgg16_model = vgg16.VGG16()
```

#--- Display architecture

```
vgg16_model.summary(show_trainable = True)
```

#--- Load the pre-trained model, VGG16, without head

```
vgg16_model = vgg16.VGG16(include_top = False)
```

```
vgg16_model.summary(show_trainable = True)
```

Code for Training, Predicting and Evaluating

#--- Train a model

```
model.compile(loss_function, metric_list)
```

```
model.fit(trainX, trainY, epochs, batch_size, validation)
```

#--- Predict by a model

```
predictY = np.argmax(model.predict(testY))
```

#--- Evaluate by model


```
model.compile(loss_function, metric_list)
```

```
loss, metric1, metric2, ... = model.evaluate(testX, testY)
```

Fully Connected Neural Network (FCNN)

FCNN consists of a series of fully connected layers

It is also known as multi-layer perceptron (MLP).




Fully Connected Layer


A fully connected layer connects every neuron in one layer to every neuron in its previous layer.

An image need to be turned into a vector before feeding into an fully connected layer.

Flatten() is used in Tensorflow.keras




CNN Classifier



Convolution in Image Processing

- In image processing,
 - Convolution is used to modify the spatial frequency characteristics of an image.




- It is a filter effect on images
- It is the process of adding each element of the image to its local neighbors, weighted by the kernel/filter.

Convolution for Single-Channel Input & Single-Channel Output

Convolution layer adds each element of a feature map to its local neighbors, weighted by the kernel.

Source pixel: 1


Destination pixel: $4 \times 0 + 0 \times 0 + 0 \times 0 + 0 \times 0 + 1 \times 0 + 2 \times 0 + 0 \times 0 + 1 \times 0 + (-4) \times 2 = -8$



Total number of multiplication for one pixel = kernel_height x kernel_width


Convolution for Multi-Channels Input & Single-Channel Output

- Number of Kernel's Channels = Number of Input Feature Map's Channels




Convolution for Multi-Channels Input & Multi-Channels Output

- Number of Kernels = Number of Output Feature Map's Channels




Padding

Padding preserves the spatial size of the input so that the output after convolution remains the same size or desired size.



Input Image

Applying padding
of 1 on 3X3



Padded Image

